

Three pivots, one **constant**.

From the outside, three pivots in six weeks looks like thrash. From the inside it was a search converging on one thing. Every version of this project—under three different names, aimed at three different tracks—was really about the same asset: **durable, manipulation-resistant agent memory**. Each pivot dropped a wrong assumption and kept that.

ORIGIN · ANIMA SWARM

An agent swarm in search of a problem

We started with **Anima Swarm**: an autonomous multi-agent system on Sui, where agents ran strategies, remembered what they learned in **MemWal**, and persisted that memory to **Walrus**. The machinery worked. What it lacked was a problem worth it—a domain with money on the table and adversaries on the other side. The agent-memory stack, the Sui deploy pipeline, and the dapp-kit frontend were all born here.

PIVOT 1 · SLABCLAW PREDICT, ON DEEPBOOK

Point the swarm at a real market

SlabClaw already had what Anima lacked: a production **10-platform price oracle over 5,167 graded cards**—a live \$2B+/yr market with real stakes. We pointed the swarm at it and built prediction markets —"Polymarket for collectibles." The obvious home looked like **DeepBook Predict**, Sui's newest DeFi primitive and the one judges most wanted to see adopted. We wrote the markets, the oracle bridge, the frontend, four live markets.

THE WALL · DEEPBOOK CAN'T PRICE A CARD

A volatility surface for an asset that has none

Then we read DeepBook Predict closely. It's a **volatility-surface options protocol**—OracleSVI, sub-hour rolling expiries, a continuous vol surface priced per strike. That exists for a liquid asset like BTC. A graded card is discrete, illiquid, and has no vol surface and no sub-hour cycle, so the protocol **structurally cannot price it**. The track's hard requirement was to integrate the DeepBook Predict contract on testnet. For our asset class that was impossible. Forcing it would have meant disqualification.

PIVOT 2 · WALRUS, THE REAL THESIS

The venue was never the innovation

The wall made the obvious thing visible: the market venue was never the point. The innovation was the **memory-backed, manipulation-resistant oracle swarm**—which is exactly what the Walrus track asks for: agents with long-term memory (MemWal), long-running stateful workflows, evidence as verifiable artifacts. So we reframed. The prediction market is the **showcase**; the agentic oracle is the **product**. And because nothing we'd built was DeepBook-specific, the market, oracle, bridge, and frontend carried over whole—**zero rework**.

The thing all three versions were circling

And here it is. **13 source-specialist agents across 11 independent venue families**, each remembering its platform's history in MemWal. A coordinator that rejects manipulation with **MAD outlier removal, anchor gates, and reputation that learns** who lied to it. **Evidence and memory snapshots on Walrus**, with the latest memory pointer **anchored onchain**. Settlement math that's **formally verified by the Sui Prover**—proven solvent for every input, not just the cases a test happens to try. Kill either production node and the swarm rebuilds its accumulated knowledge from chain and Walrus alone.

WHAT CARRIED ACROSS EACH PIVOT—NOTHING ESSENTIAL WAS THROWN AWAY

Anima Swarm	MemWal agent memory · Sui deploy pipeline · dapp-kit frontend · Walrus persistence —the memory stack
SlabClaw (app)	10-platform price oracle · 5,167-card registry with cross-grader normalization —the real market
The DeepBook build	Custom market · oracle · bridge · 4 markets · React dapp —100% reused, none was DeepBook-specific

We never pivoted away from anything that mattered. We pivoted toward the one thing every version had in common—memory that learns who tried to manipulate it, and survives being switched off. The collectibles markets are how you see it work. The memory is the product.