

# Three pivots, one **constant**.

From the outside, three pivots in six weeks looks like thrash. From the inside it was a search converging on one thing. Every version of this project — under three different names, aimed at three different tracks — was really about the same asset: **durable, manipulation-resistant agent memory**. Each pivot dropped a wrong assumption and kept that.

## ORIGIN · ANIMA SWARM

### An agent swarm in search of a problem

We started with **Anima Swarm**: an autonomous multi-agent system on Sui, where agents ran strategies, remembered what they learned in **MemWal**, and persisted that memory to **Walrus**. The machinery worked. What it lacked was a problem worth it — a domain with real money and real adversaries on the other side. The agent-memory stack, the Sui deploy pipeline, and the dapp-kit frontend were all born here.

## PIVOT 1 · SLABCLAW PREDICT, ON DEEPBOOK

### Point the swarm at a real market

SlabClaw already had what Anima lacked: a production **10-platform price oracle over 5,167 graded cards** — a real \$2B+/yr market with real stakes. We pointed the swarm at it and built prediction markets — "Polymarket for collectibles." The obvious home looked like **DeepBook Predict**, Sui's newest DeFi primitive and the one judges most wanted to see adopted. We wrote the markets, the oracle bridge, the frontend, four live markets.

## THE WALL · DEEPBOOK CAN'T PRICE A CARD

### A volatility surface for an asset that has none

Then we read DeepBook Predict closely. It's a **volatility-surface options protocol** — OracleSVI, sub-hour rolling expiries, a continuous vol surface priced per strike. That exists for a liquid asset like BTC. A graded card is discrete, illiquid, and has no vol surface and no sub-hour cycle, so the protocol **structurally cannot price it**. The track's hard requirement was to integrate the DeepBook Predict contract on testnet. For our asset class that was impossible — forcing it would have meant disqualification.

## PIVOT 2 · WALRUS, THE REAL THESIS

### The venue was never the innovation

The wall made the obvious thing visible: the market venue was never the point. The innovation was the **memory-backed, manipulation-resistant oracle swarm** — which is exactly what the Walrus track asks for: agents with long-term memory (MemWal), long-running stateful workflows, evidence as verifiable artifacts. So we reframed. The prediction market is the **showcase**; the agentic oracle is the **product**. And because nothing we'd built was DeepBook-specific, the market, oracle, bridge, and frontend carried over whole — **zero rework**.

## The thing all three versions were circling

What shipped is what every earlier version was reaching for. **13 source-specialist agents across 11 independent venue families**, each remembering its platform's history in MemWal. A coordinator that rejects manipulation with **MAD outlier removal, anchor gates, and reputation that learns** who lied to it. **Evidence and memory snapshots on Walrus**, with the latest memory pointer **anchored onchain**. Settlement math that's **formally verified by the Sui Prover** — proven solvent for every input, not just the cases a test happens to try. Kill either production node and the swarm rebuilds its accumulated knowledge from chain and Walrus alone.

WHAT CARRIED ACROSS EACH PIVOT — NOTHING ESSENTIAL WAS THROWN AWAY

### Anima Swarm

MemWal agent memory · Sui deploy pipeline · dapp-kit frontend · Walrus persistence — **the memory stack**

### SlabClaw (app)

10-platform price oracle · 5,167-card registry with cross-grader normalization — **the real market**

### The DeepBook build

Custom market · oracle · bridge · 4 markets · React dapp — **100% reused, none was DeepBook-specific**

**We never pivoted away from anything that mattered.** We pivoted toward the one thing every version had in common — memory that learns who tried to manipulate it, and survives being switched off. The collectibles markets are how you see it work. The memory is the product.